

Lecture 5: SQL's Data Definition Language

CS1106/CS6503– Introduction to Relational Databases

Dr Kieran T. Herley

2019-2020

School of Computer Science & Information Technology
University College Cork

Data Definition

Data Definition and Data Manipulation

- So far we've looked at SQL's *data manipulation language* (DML)
 - Extract information from table
 - Manipulate table by adding/removing rows of changing values
- How to we specify and create a table in the first place?
- SQL's also includes a *data definition language* (DDL) that allows us to
 - Specify the structure of your database's table(s)
 - Create a table with this structure

A Table Definition for Our Running Example



```
CREATE TABLE students
(
  id_number CHAR(9),
  first_name VARCHAR(20),
  last_name VARCHAR(30),
  date_of_birth DATE,
  hometown VARCHAR(30),
  course CHAR(5),
  points INTEGER,
  . . .
);
```

- Creates table named `students` with specified structure
- Structure
 - specified as list of attribute names and associated *types*
 - judgement required in choosing appropriate types (and lengths)
- (Health warning: Above could be improved!)

Overview of SQL's Data Types

- Remember that databases are used in a huge variety of applications to house very diverse types of data
- SQL provides a range of different “data types” with which to populate our tables

Textual Strings of characters

Numerical Numbers both integers and real

Temporal Dates, times, date-time timestamps

Others

- Also other types: BLOB (Binary Large Object)
- Most DBMSs support many flavours; often system dependent

Integer Types

- For integer (whole) numbers we use INTEGER (or INT) e.g.

num_points **INTEGER**

- Range of allowed values (MySQL ¹)

$-2,147,483,468 \dots 2,147,483,467$

MySQL should be good enough for most purposes.

- Most systems offer variants (SMALLINT, BIGINT) with different ranges; require different amount of space; need to choose carefully

¹Limits typically DBMS specific.

Other Numerical Types

DECIMAL(*n*, *d*)

- any *n*-digit number with a *d*-digit mantissa
-

temperature **DECIMAL**(4, 1)

123.4 ok; 12.34 rounded to 12.3; 1234.5 error

FLOAT

- scientific notation (e.g. 1.34E+12 for $1.34 \times 10^{+12}$)
- Useful for scientific data
- system-dependent limits on number size

DATE

- Dates in YYYY-MM-DD format

TIME

- Time (24-hour clock) in hh:mm:ss format

Others e.g. DATETIME **MySQL**

NOTES DBMSs support useful *functions* for manipulating dates and times e.g. MONTH(.), DATE_ADD(.,.) **MySQL**

CHAR(*n*)

- Shortish, fixed-length strings
 - Space for exactly *n* characters allocated
 - shorter strings right-padded with blanks, longer ones truncated
- Useful for id numbers, course codes etc.

VARCHAR(*n*)

- Strings of any any length up to max of *n* characters
- May be more space-efficient than CHAR
- Useful where precise string length not fixed or known in advance, e.g. addresses

Note Both draw individual characters from underlying character set—`latin1_swedish_ci` by default in our case

- Need to be careful in choosing appropriate “size”
- n too large: wasteful of space (CHAR)
- n too small: String may be truncated if assigned to a CHAR(n)/VARCHAR(n) column of insufficient “size” i.e only first n characters retained

Meanwhile Back At The Example



```
CREATE TABLE students
(
    id_number CHAR(9),
    first_name VARCHAR(20),
    last_name VARCHAR(30),
    date_of_birth DATE,
    hometown VARCHAR(30),
    course CHAR(5),
    points INTEGER,
    . . .
);
```

- Should only execute CREATE once when table is first set up
- Can also specify other useful information within CREATE statement: Information about DB keys, Default values for columns, Constraints on values. Will look at this later.

Other Structure-Altering SQL Commands

- To expunge a table:

```
DROP TABLE X;  
DROP TABLE IF EXISTS Y;
```

- Careful– deletes table and contents
- Altering table structure
 - Adding an attribute/column:

```
ALTER TABLE students ADD gender CHAR(1);
```

- Deleting an attribute/column:

```
ALTER TABLE students DROP hometown;
```

If you design your DB properly, you should rarely need these

students

<i>id_number</i>	<i>first_name</i>	<i>last_name</i>	<i>date_of_birth</i>	<i>hometown</i>	<i>course</i>	<i>points</i>
112345678	Aoife	Ahern	1993-01-25	Cork	ck401	500
112467389	Barry	Barry	1980-06-30	Tralee	ck402	450
112356489	Ciara	Callaghan	1993-03-14	Limerick	ck401	425
112986347	Declan	Duffy	1993-11-03	Cork	ck407	550
112561728	Eimear	Early	1993-07-18	Thurles	ck406	475
112836467	Fionn	Fitzgerald	1994-06-13	Bandon	ck405	485

- Each table should have one or more attributes (collectively known as the *key*) the values of which uniquely identify each row i.e. no two rows should have the same key
- Example: `id_number` plays this role in `students`
- Table definition should specify key as shown

```
CREATE TABLE students
(
    . . .
    PRIMARY KEY (id_number)
)
```

- Yes, there are also non-primary keys; more on this later

Table Specification

```
CREATE TABLE guests
(  id_number INTEGER AUTO.INCREMENT, -- new rows automatically issued fresh id
                                     -- use with primary key column only
   name VARCHAR(20) NOT NULL,        -- inserting/setting NULL will trigger error
   nationality VARCHAR(20) DEFAULT 'Irish' -- all guests Irish unless specified otehrwise

   PRIMARY KEY (id_number)
);
```

Table Specification

```
CREATE TABLE guests
( id_number INTEGER AUTO.INCREMENT, -- new rows automatically issued fresh id
                                   -- use with primary key column only
  name VARCHAR(20) NOT NULL,      -- inserting/setting NULL will trigger error
  nationality VARCHAR(20) DEFAULT 'Irish' -- all guests Irish unless specified otehrwise

  PRIMARY KEY (id_number)
);
```

Some Insertions

```
INSERT INTO guests (name, nationality) VALUES
( 'Paddy', ' Irish ' ),    ( 'Gunther', 'German'),    ( ' Luigi ', ' Italian ' );
INSERT INTO guests (name) VALUES
( 'Seamus' );
INSERT INTO guests (name, nationality) VALUES
( NULL, 'Irish ' );        -- ERROR!!!
```

Creating and Populating A Database

- cs1106 website houses two files used to create clones of students DB in your accounts
 - `students_setup.sql`: file containing CREATE statement
 - `students_populate.sql`: file containing a bunch of INSERT statements
- MySQL can process SQL commands taken from a file not just those type in; in HeidiSQL use the folder icon to the right of the blue execution button
- We will provide files for most of the sample DBs we use in this module

An Example

A Simple Database

- Suppose we want to design a DB to hold information about some people



Information about each person:

- Name
- Birth date
- Address
- Favourite foods

A Simple Database

- Suppose we want to design a DB to hold information about some people

-

Information about each person - First stab at a DB design:
son:

	Column	Type
• Name	name	VARCHAR(?)
• Birth date	gender	CHAR(1)
	birth_date	DATE
• Address	address	VARCHAR(?)
• Favourite foods	favourite_foods	?????

A Simple Database

- Suppose we want to design a DB to hold information about some people

-

Information about each person
son: First stab at a DB design:

	Column	Type
• Name	name	VARCHAR(?)
• Birth date	gender	CHAR(1)
	birth_date	DATE
• Address	address	VARCHAR(?)
• Favourite foods	favourite_foods	?????

- Some imperfections with this
 - Best to have unique label for each person in case of duplicate names; introduce ID numbers
 - Names are really composite objects; split into first and last names; same with addresses
 - Need to choose sensible values for VARCHAR widths

A Second Stab

-

Column	Type
person_id	CHAR(6)
first_name	VARCHAR(20)
last_name	VARCHAR(20)
gender	CHAR(1)
birth_date	DATE
street	VARCHAR(30)
town	VARCHAR(30)
county	VARCHAR(30)
favourite_foods	?????

- What to do about favourite foods?
 - VARCHAR– difficult to access individual food items
 - Separate columns (fav1, fav2, . . .), but how many?
 - Better to use second separate table to capture this “relationship” (persons to food types)

The favourite_foods Table



persons			
<i>person_id</i>	<i>first_name</i>	<i>last_name</i>	...
...
112356489	Ciara	Callaghan	...
112986347	Declan	Duffy	...
...

favourite_foods	
<i>person_id</i>	<i>food</i>
...	...
...	...
112356489	Ice cream
112356489	Chocolate
112986347	Pizza
112986347	Beer
112986347	Crisps
...	...

- This models the fact that person 112986347 (aka Declan Duffy) likes pizza, beer and crisps
- The “link” between the two tables is the *person_id*; within the favourite_foods table the *person_id* is a *foreign key* that references the persons table.

Our Complete Design

```
CREATE TABLE persons
(
    person_id  CHAR(6),
    first_name VARCHAR(20),
    last_name  VARCHAR(20),
    gender     CHAR(1),
    birth_date DATE,
    street     VARCHAR(30),
    town       VARCHAR(30),
    county     VARCHAR(30),

    PRIMARY KEY (person_id)
);
```

```
CREATE TABLE favourite_foods
(
    person_id CHAR(6),
    food      VARCHAR(20),

    PRIMARY KEY (person_id, food)
);
```

Note: two-attribute key

MySQL enforces key-distinctness property for tables, e.g. disallow insertion of duplicate person_id values into persons

Working With Our New DB

- Can pose queries at either table using SQL tools we've learnt

```
SELECT *  
FROM persons  
WHERE first_name = 'Ciara';
```

```
SELECT *  
FROM favourite_foods  
WHERE person_id = '112986347';
```

- What about queries like the following?

List names of all persons who like pizza

What We Have Covered So Far

- Setting up a simple database (CREATE)
- Adding content to the database (INSERT, UPDATE)
- Posing (simple) queries to extract information from database

The favourite foods example is taken from “Learning SQL” by Alan Beaulieu (O'Reilly, 2009). If you are looking for a nice, compact and affordable introduction to SQL, this is a good choice.