

# Lecture 6: Full-Relation Queries

*CS1106/CS6503– Introduction to Relational Databases*

---

Dr Kieran T. Herley

2019-2020

School of Computer Science & Information Technology  
University College Cork

## Summary

*SQL's aggregation functions. Aggregation and grouping.*

# Aggregation and Grouping

- Our queries to data all relate to individual rows within tables (possibly more than one), not queries that “summarize” data from the entire table (or sections thereof)
- Today’s DB contains stats on the countries of the world (population, GDP etc.)
- How to answer questions like:
  - What is the population of Europe?
  - What country has the largest GDP?
  - Which region in the world has the greatest area?

# Today's Database

## countries

<i>name</i>	<i>region</i>	<i>area</i>	<i>population</i>	<i>gdp</i>
Afghanistan	South Asia	652225	26000000	NULL
Albania	Europe	28728	3200000	6656000000
Algeria	Middle East	2400000	32900000	75012000000
:	:	:	:	:
Ireland	Europe	70182	4000000	137120000000
:	:	:	:	:
Zimbabwe	Africa	390759	12900000	6192000000

## A Closer Look At countries

countries				
<i>name</i>	<i>region</i>	<i>area</i>	<i>population</i>	<i>gdp</i>
Afghanistan	South Asia	652225	26000000	NULL
Albania	Europe	28728	3200000	6656000000
Algeria	Middle East	2400000	32900000	75012000000
.	.	.	.	.
.	.	.	.	.
.	.	.	.	.
Ireland	Europe	70182	4000000	137120000000
.	.	.	.	.
.	.	.	.	.
Zimbabwe	Africa	390759	12900000	6192000000

1

**name** countries name (as VARCHAR)

**region** surrounding region (as VARCHAR)

**area** in km<sup>2</sup> (as DECIMAL )

**population** (as DECIMAL)

**gdp** Gross Domestic Product in dollars; (as DECIMAL )

<sup>1</sup>Data may be somewhat “stale”, but will suffice for target practice.

# What Is The Population Of China?

- China's population:

```
SELECT name, population  
FROM countries  
WHERE name = 'China';
```

# What Is The Population Of China?

- China's population:

```
SELECT name, population  
FROM countries  
WHERE name = 'China';
```

- What does this do?

```
SELECT name, area/population  
FROM countries  
WHERE region = 'Europe'  
ORDER BY area/population DESC;
```

Note the use of *expressions* like `area/population` in `SELECT` and `ORDER BY` clauses

# Whole-Table Queries Using Aggregates

---



## How Many Countries Are There?

- Desired result summarizes info. from entire table not just one or two rows.

# How Many Countries Are There?

- Desired result summarizes info. from entire table not just one or two rows.

- 

```
SELECT COUNT(name)  
FROM countries;
```


COUNT(name)
<hr/>
189

- Counts the number of values in the `name` column
- Result (shown right) is a  $1 \times 1$  table!
- Could use `COUNT(*)` instead of `COUNT(name)` here

## How Many Countries Are There cont'd

- 

```
SELECT SUM(area)
FROM countries;
```



- SQL has five main *aggregation* operators:
  - SUM** sum of column values
  - AVG** average of column values
  - MIN** minimum of column values
  - MAX** maximum of column values
  - COUNT** count of values in column with duplicates counted every time they appear
- All these can be used to compute aggregates of *groups* of rows sharing some common characteristics

## Some Whole-Table Aggregation Queries

- What is the total population of the world?

```
SELECT SUM (population)  
FROM countries;
```

## Some Whole-Table Aggregation Queries

- What is the total population of the world?

```
SELECT SUM (population)  
FROM countries;
```

- What is the largest GDP?

```
SELECT MAX (gdp)  
FROM countries;
```

## Some Whole-Table Aggregation Queries

- What is the total population of the world?

```
SELECT SUM (population)  
FROM countries;
```

- What is the largest GDP?

```
SELECT MAX (gdp)  
FROM countries;
```

## Some Whole-Table Aggregation Queries

- How many regions are there?

```
SELECT COUNT(region)  
FROM countries;
```

Wrong!

## Some Whole-Table Aggregation Queries

- How many regions are there?

```
SELECT COUNT(region)  
FROM countries;
```

**Wrong!** counts the num. values in `region` column including duplicates

- Second go at “How many regions are there?”

```
SELECT COUNT(DISTINCT region)  
FROM countries;
```

- `DISTINCT` ensures each region is counted only once
- Can use this with any aggregate operator but only makes sense with `COUNT`



# GROUPING

---

# What is the Breakdown of Population By Region?

- Can specify grouping of rows using GROUP BY and then the aggregates apply to each group separately

- 

```
SELECT SUM(population)  
FROM countries  
GROUP BY region;
```

SUM(population)

668981725

7432900

2173791900

818020000

373882000

433400000

358510000

1462138000

- Conceptually this works as follows:
  - Rows with same region value are bunched together into groups
  - SUM applied to each group separately
  - One row per group appears in result
- When we omit the GROUP BY clause, default group (the entire table) is used for any aggregate operators

# What is the Breakdown cont'd

```
SELECT SUM(population)
FROM countries
GROUP BY region;
```

## countries

<i>name</i>	<i>...</i>	<i>population</i>	<i>...</i>
Afghanistan		26000000	
Albania	...	3200000	...
Algeria	...	32900000	...
.	.	.	.
.	.	.	.
.	.	.	.
Ireland	...	4000000	...
.	.	.	.
.	.	.	.
.	.	.	.
Zimbabwe	...	12900000	...

The original table

Notes: the original table remains unaltered– the “grouping” is notional

# What is the Breakdown cont'd

```
SELECT SUM(population)
FROM countries
GROUP BY region;
```

<i>name</i>	<i>...</i>	<i>population</i>	<i>...</i>
Angola	...	14500000	...
.	.	.	.
.	.	.	.
.	.	.	.
Zimbabwe	...	12900000	...
Antigua	...	77000	...
.	.	.	.
.	.	.	.
.	.	.	.
St Vincent	...	121000	...
.	.	.	.
.	.	.	.
.	.	.	.

Grouped by region

Notes: the original table remains unaltered– the “grouping” is notional

## What is the Breakdown cont'd

```
SELECT SUM(population)
FROM countries
GROUP BY region;
```

SUM(pop)

668981725  
7432900  
2173791900  
818020000  
433400000  
358510000  
1462138000

Result containing

one row

per group,

e.g. sum

## What is the Breakdown cont'd

- Can include the grouping attribute in SELECT clause for readability

- 

```
SELECT region, SUM(population)
FROM countries
GROUP BY region;
```

region	SUM(population)
Africa	668981725
Americas	7432900
Asia-Pacific	2173791900
⋮	⋮

- SELECT clause can include only
  1. aggregate terms (e.g. SUM(. . .))
  2. un-aggregated attributes (e.g. region), but this only makes sense for attributes mentioned in GROUP BY clause

# What Country Has The Largest Population?

- First attempt:

```
SELECT name, MAX(population)  
FROM countries;
```

# What Country Has The Largest Population?

- First attempt:

```
SELECT name, MAX(population)
FROM countries;
```

Wrong!

- Some versions of SQL object (syntax error) to inappropriate mixing of aggregates and non-aggregates in SELECT clause with no GROUP columns specified
- Others includes *arbitrary* value from name column not necessarily the one corresponding to largest population:

<b>name</b>	<b>MAX(population)</b>
Albania	1300000000

- Will return to this later . . .



## Some More Examples

- List max, min and average and total GDP for each region

```
SELECT region, MAX(gdp), MIN(gdp), AVG(gdp), SUM(gdp)  
FROM countries  
GROUP BY region;
```

- List Max, min and average and total GDP for regions Europe and North America

```
SELECT region, MAX(gdp), MIN(gdp), AVG(gdp), SUM(gdp)  
FROM countries  
WHERE region IN ('Europe', 'North America')  
GROUP BY region;
```

## Some More Examples

List the regions and size (num of countries) for all regions of size  $\geq 10$

```
SELECT region, COUNT(*)  
FROM countries  
WHERE COUNT(*)  $\geq$  10  
GROUP BY region;
```

Wrong!

- WHERE conditions apply to individual table rows
- Can't use aggregates in WHERE clause

**HAVING**

---

## What Regions Contain At Least 10 Countries?

- Want to filter by some characteristic of groups rather than of individual rows, the group size in the previous example

## What Regions Contain At Least 10 Countries?

- Want to filter by some characteristic of groups rather than of individual rows, the group size in the previous example
- 

```
SELECT region, COUNT(*)  
FROM countries  
GROUP BY region;
```

Counts num. countries for all regions

```
SELECT region, COUNT(*)  
FROM countries  
GROUP BY region  
HAVING COUNT(*) >= 10;
```

Counts num. countries for regions with at least ten countries

- Aggregates in HAVING clause apply to the group in question (here region groups); The only attributes that can appear in the HAVING clause un-aggregated are those mentioned in the GROUP BY clause (i.e. region in this case).

## Which Regions Have Populations in Excess of a Billion?



```
SELECT region, SUM(population)
FROM countries
GROUP BY region
HAVING SUM(population) > 1000000000;
```

# What Is The Average GDP Of Regions With Large Populations

- List the number of countries, average GDP and total area of all regions that have a population greater than one billion
- 

```
SELECT region, COUNT(*), AVG(gdp), SUM(area)
FROM countries
GROUP BY region
HAVING SUM(population) > 1000000000;
```

## Which Regions Are Economically Homogeneous?

- List the max and min per capita GDP for all regions where the maximum GDP per capita is at most five times the the minimum GDP per capita
- 

$$\text{Per capita GDP} = \text{GDP} / \text{population}$$



# Which Regions Are Economically Homogeneous?

- List the max and min per capita GDP for all regions where the maximum GDP per capita is at most five times the the minimum GDP per capita

- 

$$\text{Per capita GDP} = \text{GDP} / \text{population}$$

- 

```
SELECT region, MIN(gdp/population), MAX(gdp/population)
FROM countries
GROUP BY region
HAVING MAX(gdp/population) < 5 * MIN(gdp/population);
```

# Which Country Has The Largest Population?



```
SELECT name, population
FROM countries
WHERE population = MAX(population);
```

**Wrong!**– Can't use aggregates in WHERE clause



```
SELECT name, population
FROM countries
HAVING population = MAX(population);
```

**Wrong!**– Can't use (un-aggregated) attributes in HAVING clause other than those in GROUP BY clause (here none)

# Which Country Has The Largest Population?

- Finally an approach that works:

```
SELECT name, population  
FROM countries  
WHERE population = (SELECT MAX(population) FROM countries)
```

- Uses *subqueries* (seen later); “inner query” (within parentheses) determines max value and “outer query” uses it to determine associated country

## Notes and Acknowledgements

The countries DB was taken from the website [www.sqlzoo.net](http://www.sqlzoo.net).  
The data itself was ultimately taken from the BBC.